

# Tree-Structured Indexes

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro

# Introduction

*As for any index, 3 alternatives for data entries  $k^*$ :*

- Data record with key value  $k$
- $\langle k, \text{rid of data record with search key value } k \rangle$
- $\langle k, \text{list of rids of data records with search key } k \rangle$

Assignment Project Exam Help

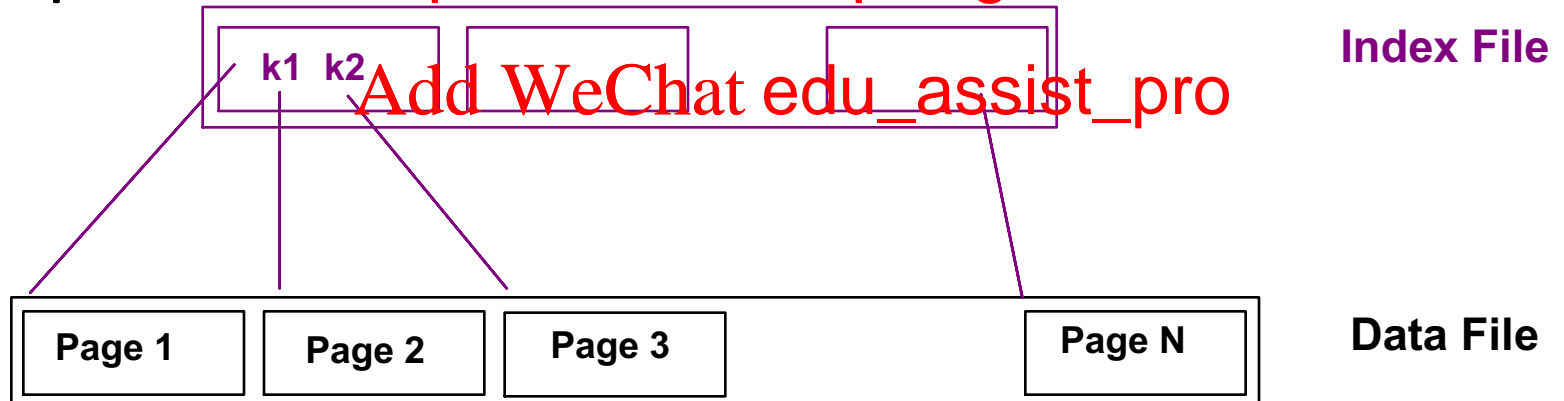
- Choice is <https://eduassistpro.github.io/> indexing technique used to locate data on  
Add WeChat edu\_assist\_pro
- Tree-structured indexing techniques support both *range searches* and *equality searches*.
- ISAM: static structure; B+ tree: dynamic, adjusts gracefully under inserts and deletes.

# Range Searches

``Find all students with  $gpa > 3.0$ ``


- If data is in sorted file, do binary search to find first such student, then scan to find others.
- Cost of binary search can be quite high.

Simple idea: <https://eduassistpro.github.io/>



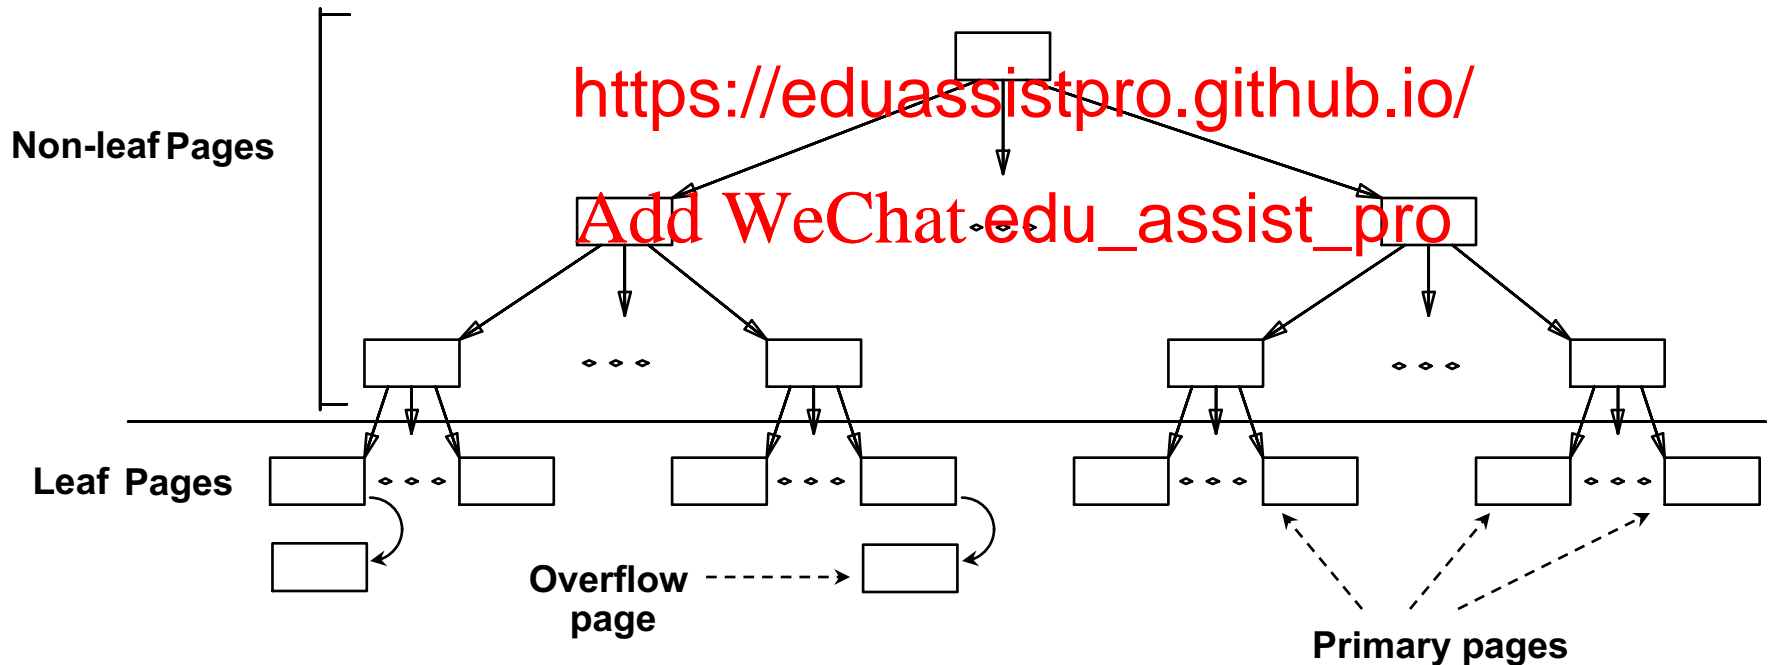
👉 Can do binary search on (smaller) index file!

# ISAM (Indexed Sequential Access Method)



Index file may still be quite large. But we can apply the idea repeatedly!

**Assignment Project Exam Help**



👉 Leaf pages contain *data entries*.

# Comments on ISAM

- *File creation*: Leaf (data) pages allocated sequentially, sorted by each key; then index pages allocated, then space for overflow pages.
- *Index entries*: <search key value, page id>; they 'direct' search for data entries, which are in leaf pages.
- Search: Start at root to go to leaf. Cost =  $\log_F N$ ;  
F = # entries/index pg, N = # leaf
- Insert: Find leaf data entry belongs to, and put it there.
- Delete: Find and remove from leaf; if empty overflow page, de-allocate.

➡ **Static tree structure**: *inserts/deletes affect only leaf pages.*

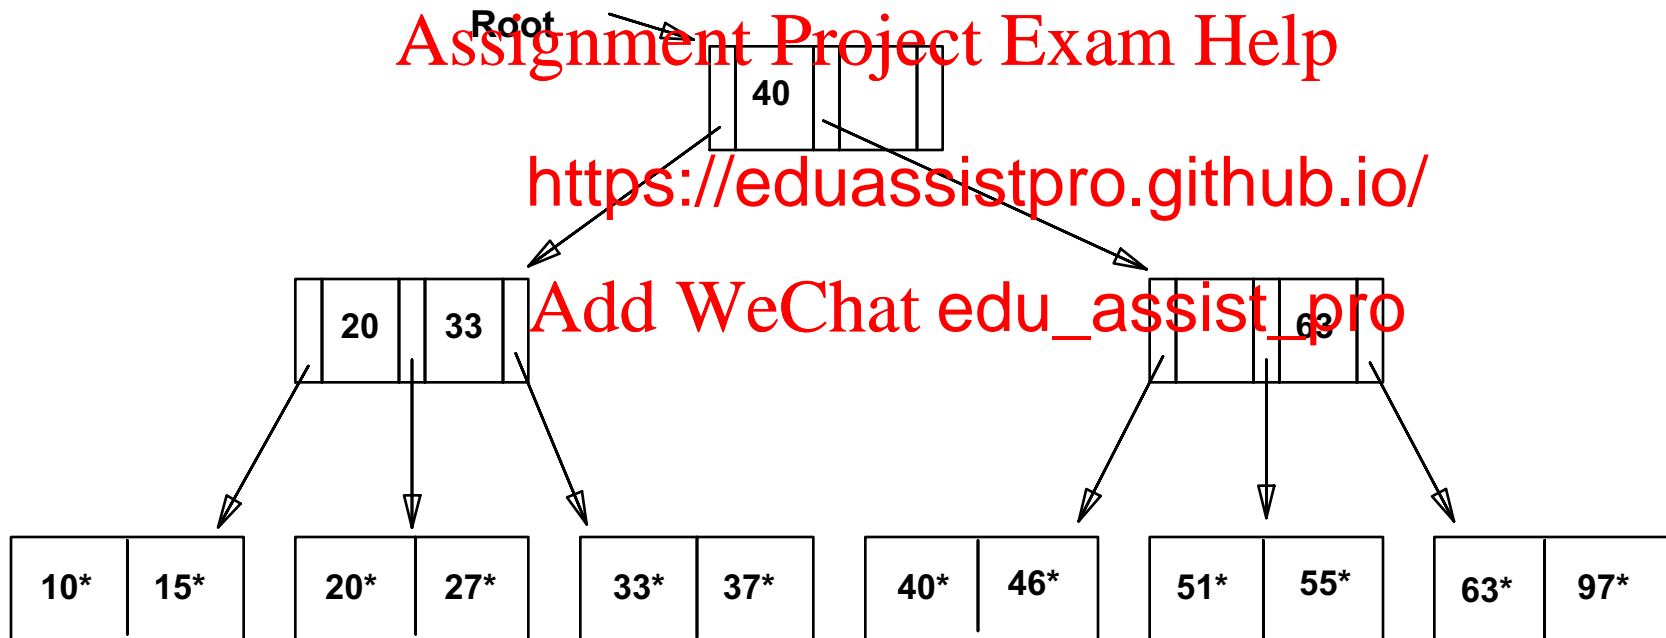
Data Pages

Index Pages

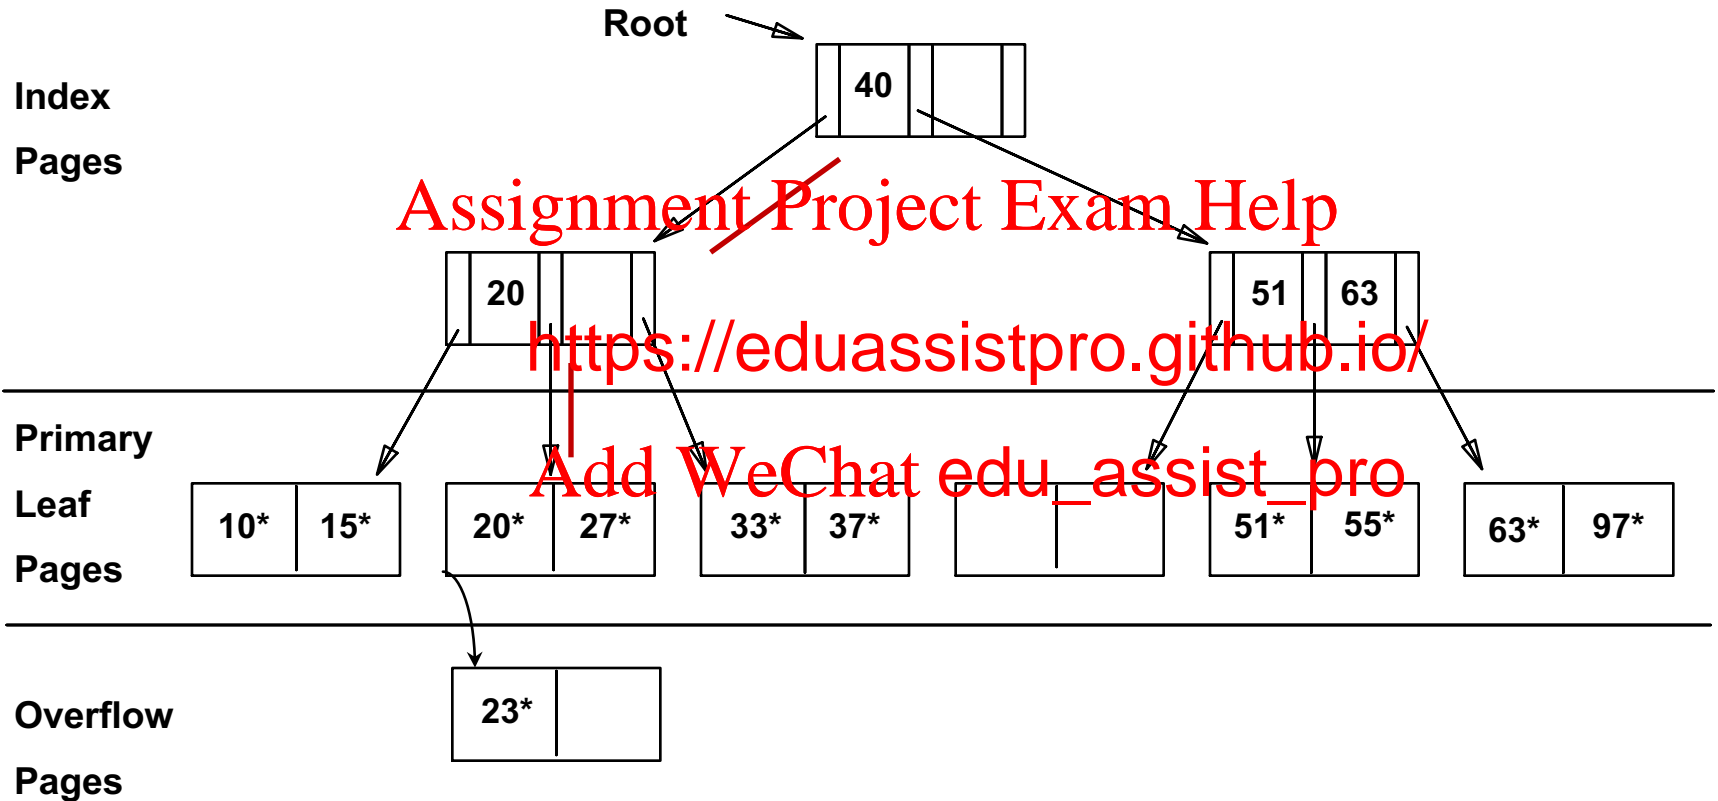
Overflow pages

# Example ISAM Tree

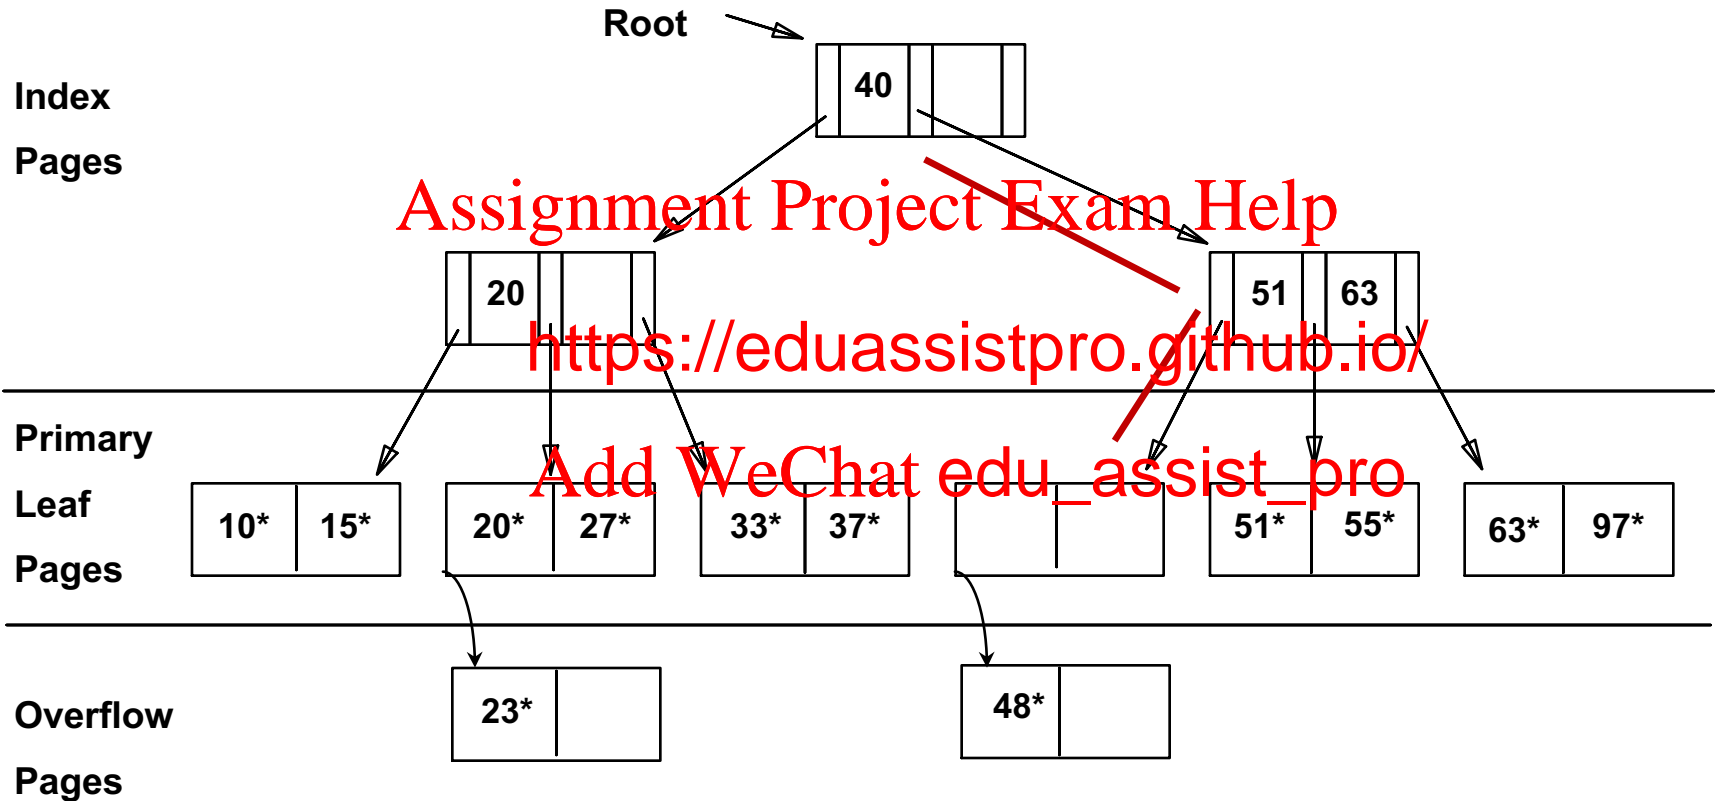
Each node can hold 2 entries; no need for 'next-leaf-page' pointers. (Why?)



# Inserting 23\*

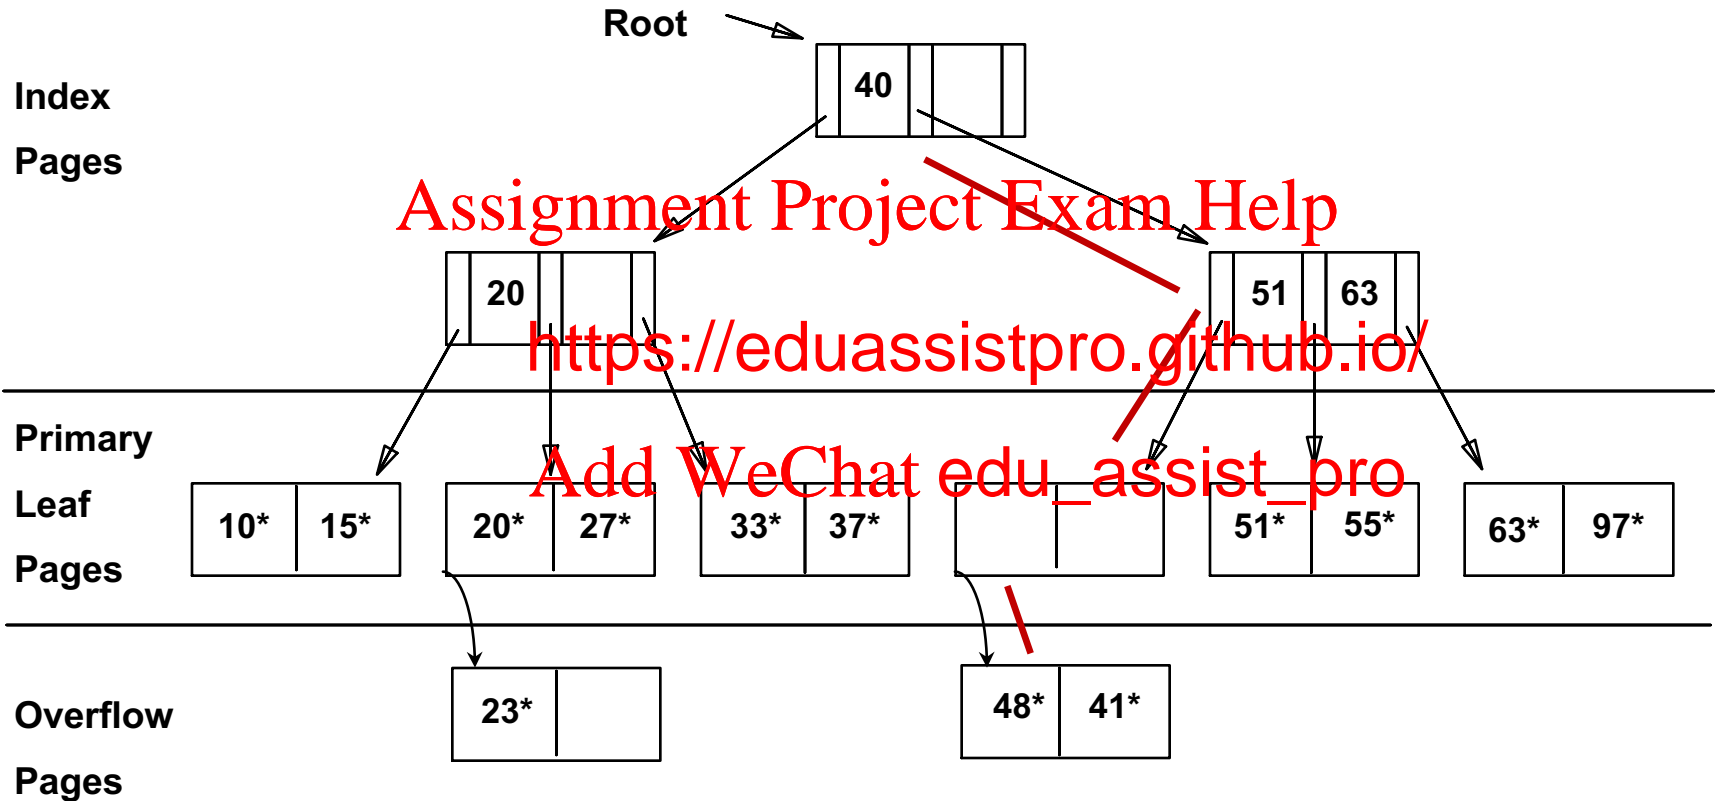


# Inserting 48\*

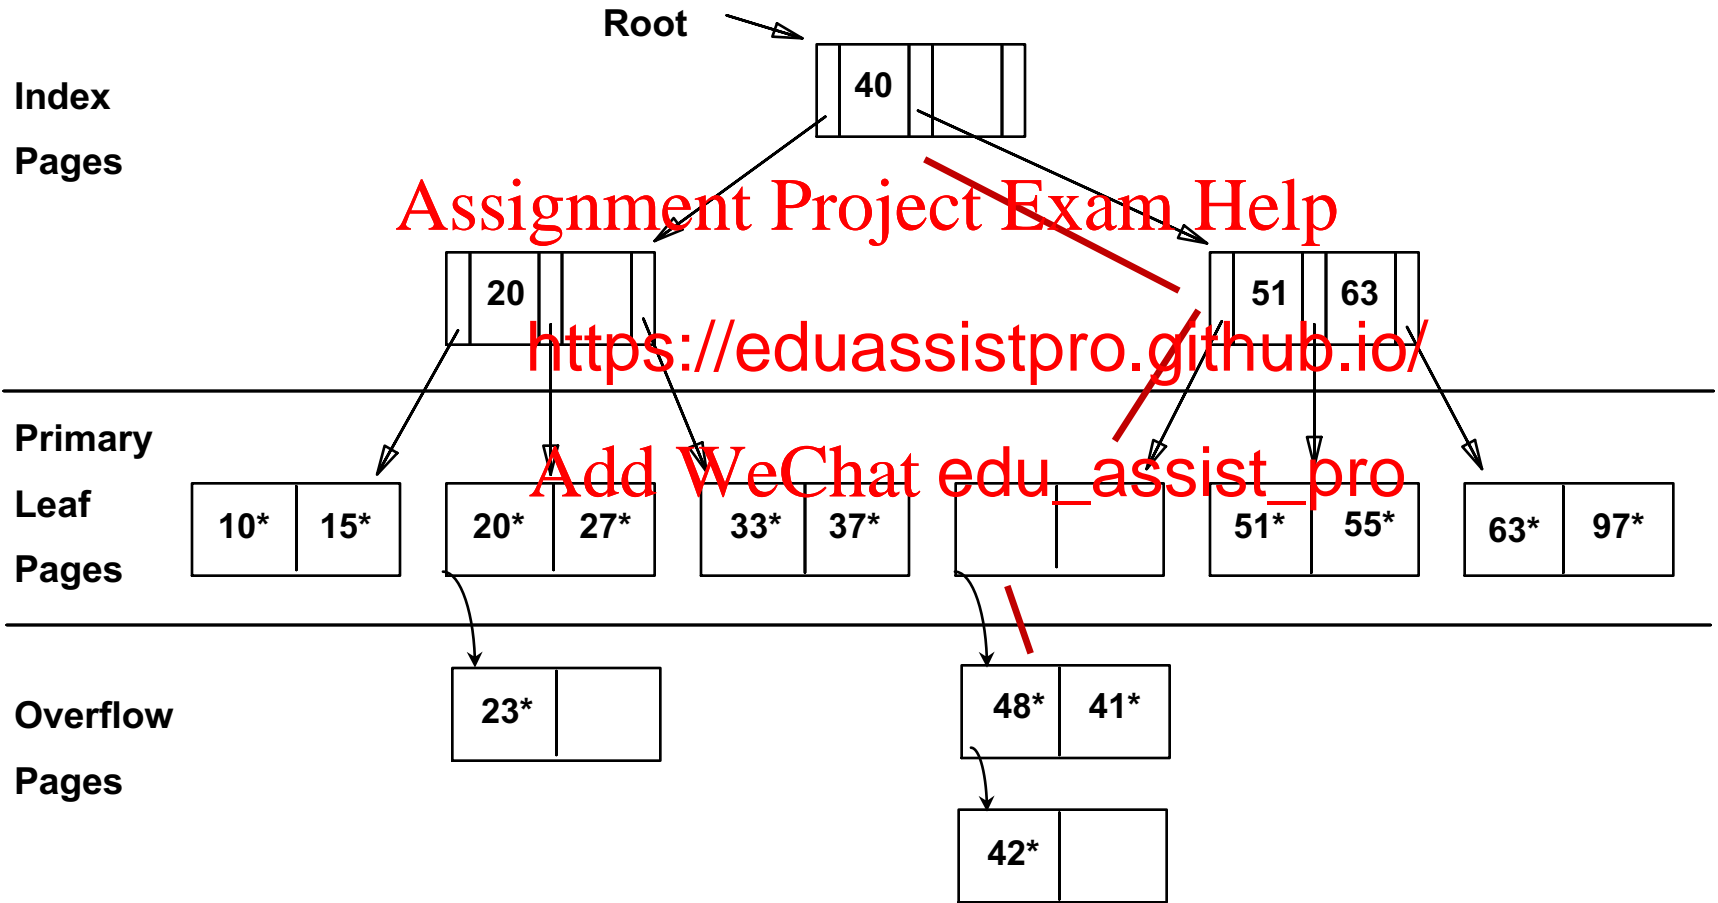




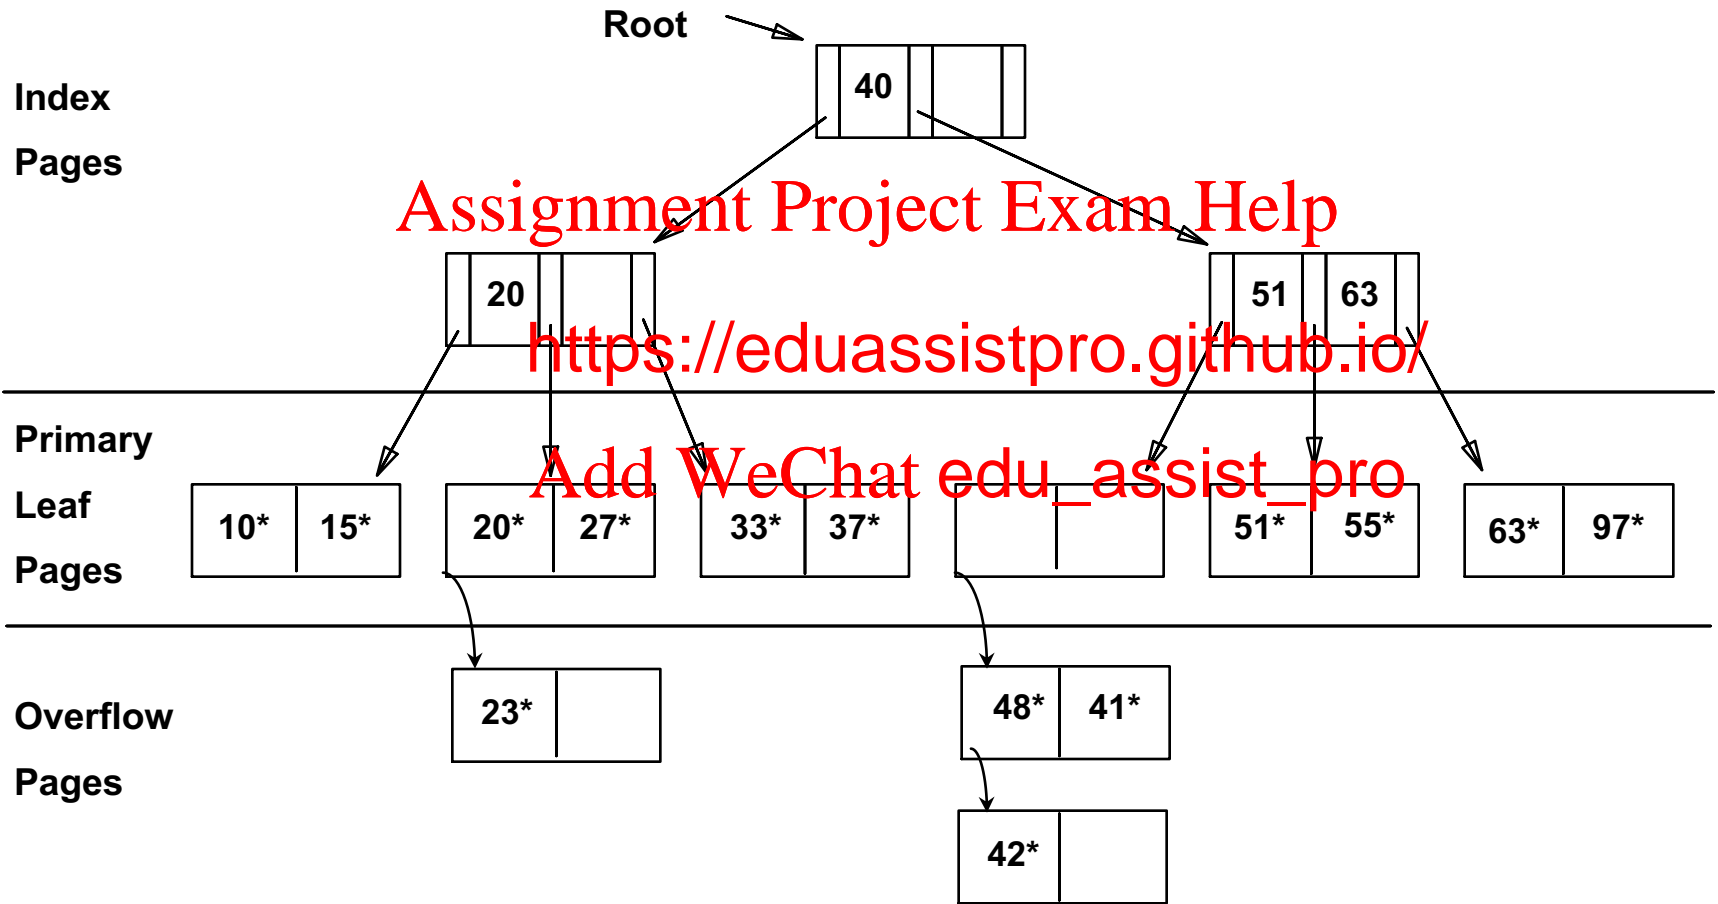
# Inserting 41\*



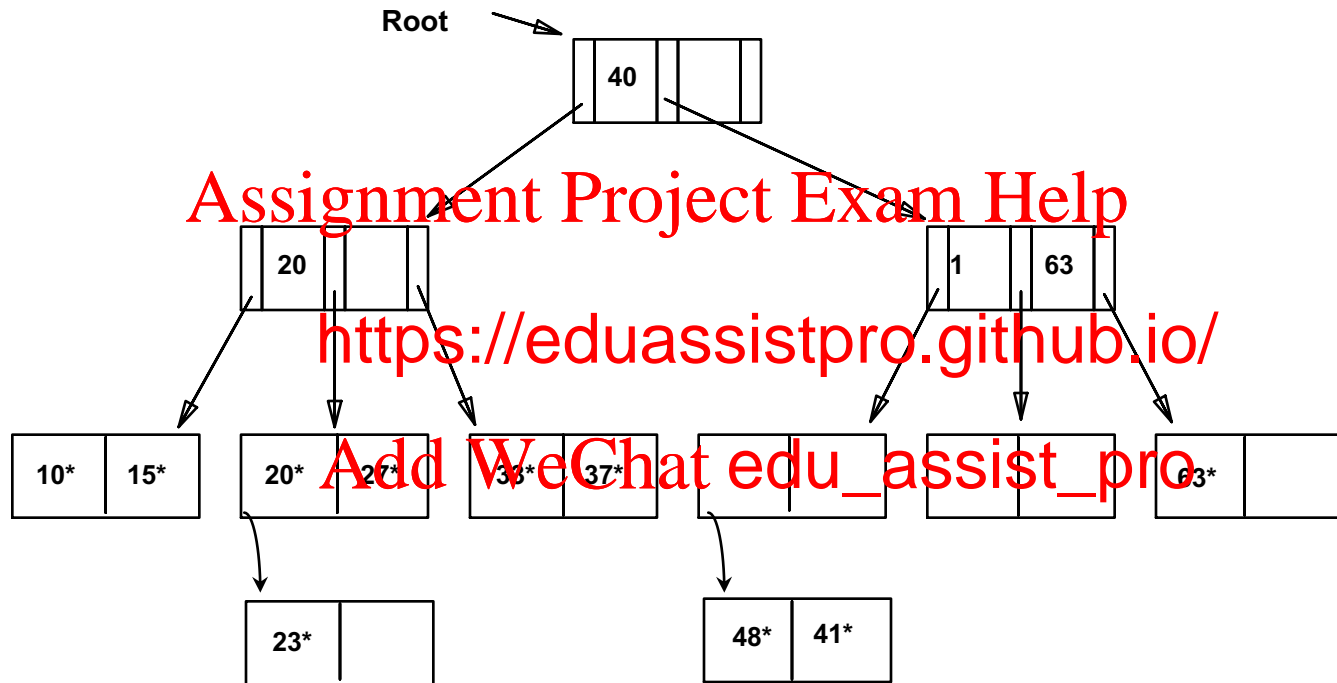
# Inserting 42\*



# Then Deleting 42\*, 51\*, 97\*, 55\*




... After Deleting 42\*, 51\*, 97\*, 55\*



👉 Note that 51\* appears in index levels, but not in leaf!

# B+ Tree: Most Widely Used Index

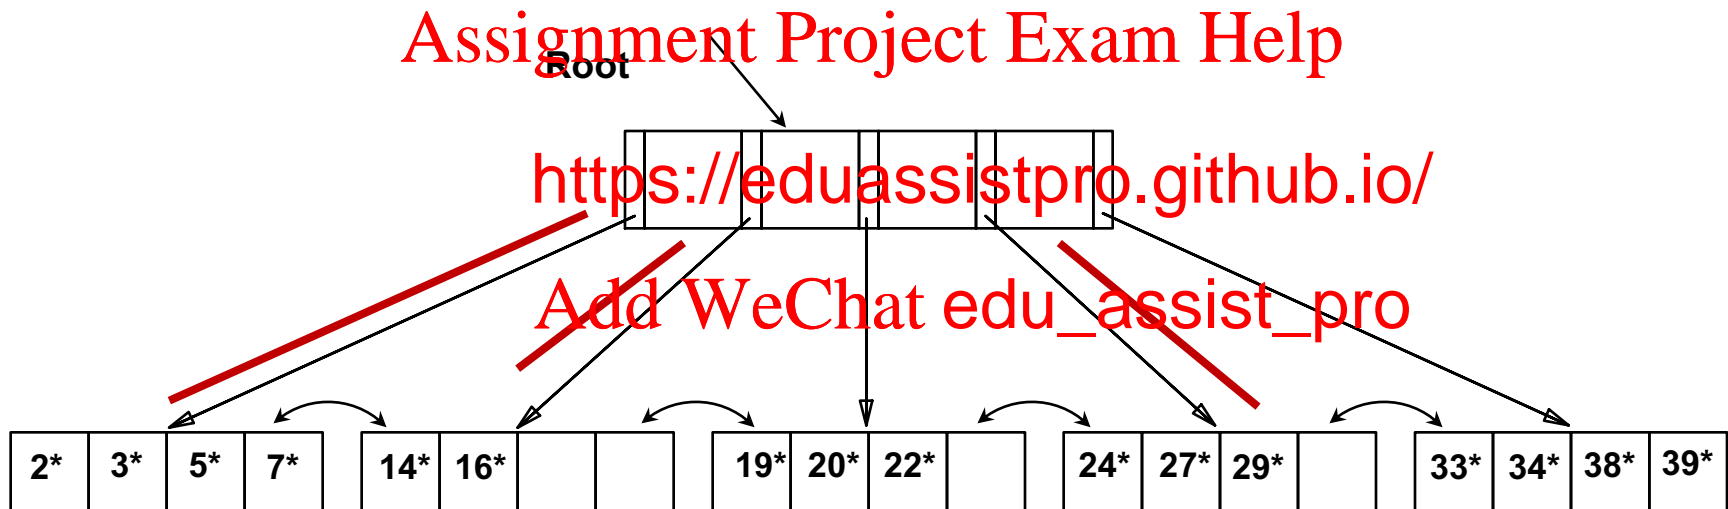
- Insert/delete at  $\log_F N$  cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- Minimum 50% occupancy (except for root). Each node contains  $d \leq \frac{m}{2}$  entries.
- The parameter  $m$  is the order of the tree.
- Supports eq



# Example B+ Tree

Search begins at root, and key comparisons direct it to a leaf.

Search for 5\*, 15\*, all data entries  $\geq 24^*$  ...



👉 *Based on the search for 15\*, we know it is not in the tree!*

# B+ Trees in Practice

Typical order: 100. Typical fill-factor: 67%.

- average fanout = 133

Typical capacities:

- Height 4:  $133^4 = 312,900,700$  records
- Height 3: 1 <https://eduassistpro.github.io/>

Can often hold top levels in pool.

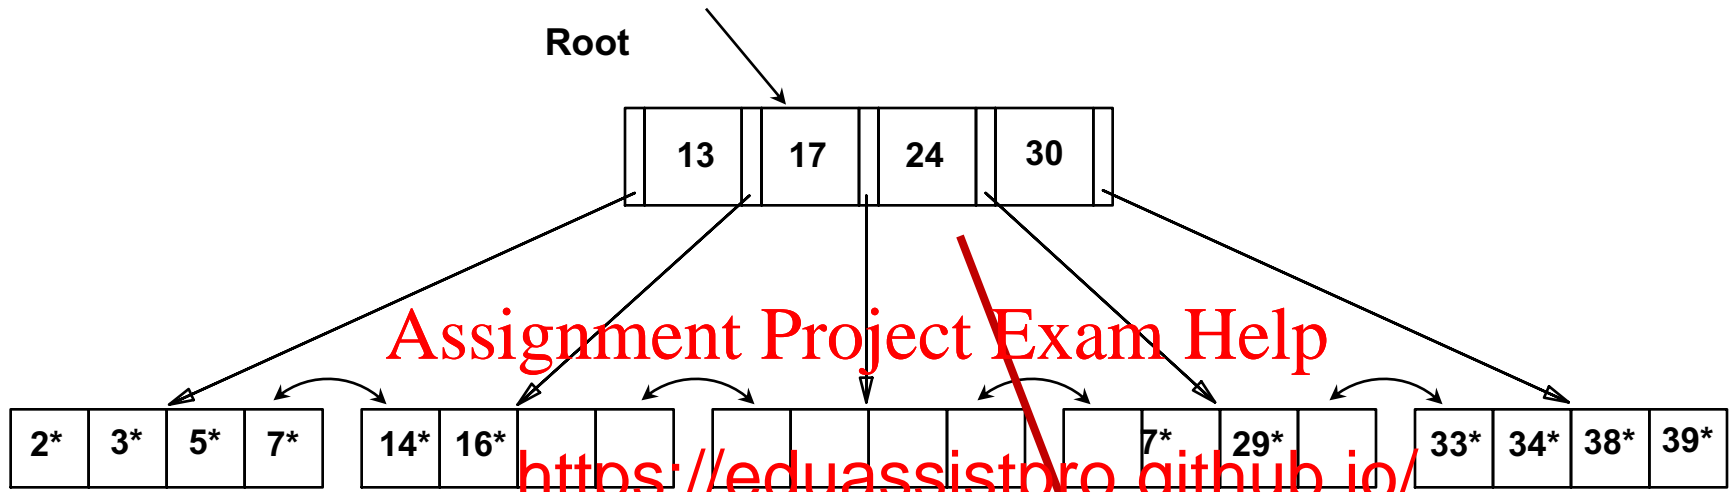
- Level 1 = 1 page = 8 Kbytes
- Level 2 = 133 pages = 1 Mbyte
- Level 3 = 17,689 pages = 133 MBytes

# Inserting a Data Entry into a B+ Tree



- Find correct leaf  $L$ .
- Put data entry onto  $L$ .
  - If  $L$  has enough space, *done!*
  - Else, must split  $L$  (node  $L2$ )
    - Redistribute middle key to  $L$  and  $L2$
    - Insert index entry pointing to  $L2$  at end of  $L$ .
- This can happen recursively
  - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- Splits “grow” tree; root split increases height.
  - Tree growth: gets wider or one level taller at top.




# Inserting 8\*



Add WeChat edu\_assist\_pro



# After Inserting 8\*



Add WeChat edu\_assist\_pro

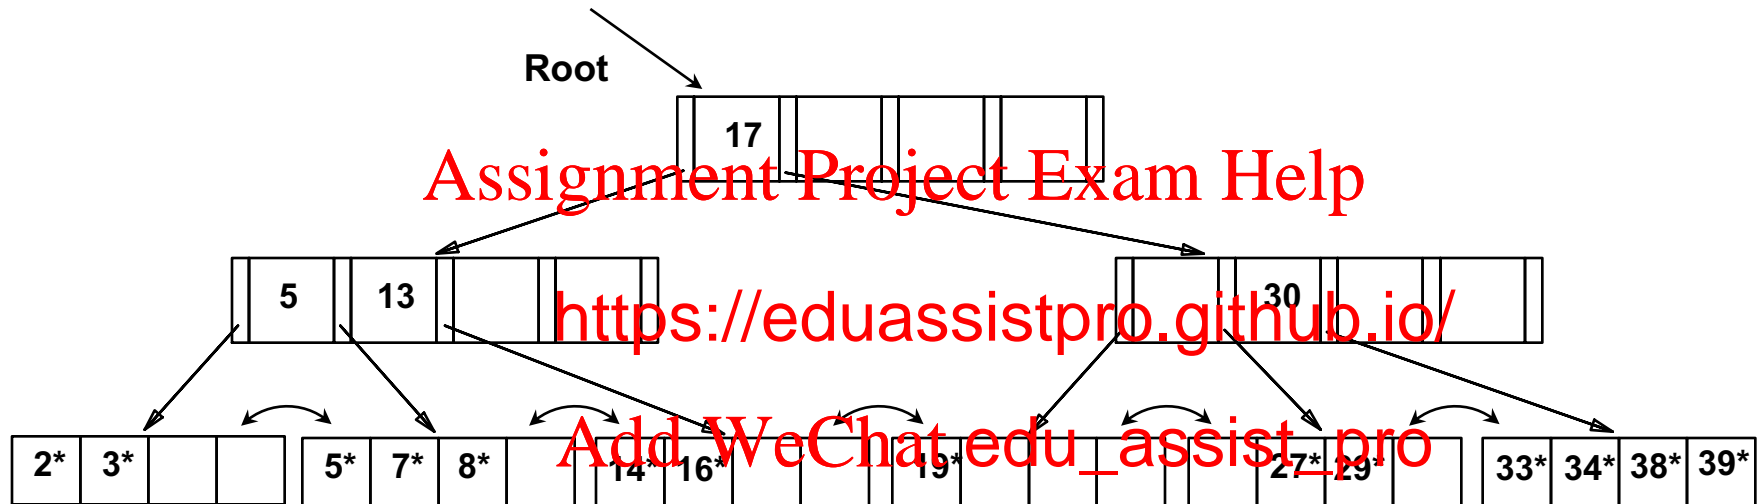
❖ In this example, we can avoid redistributing entries; however, this is usually not done in practice.

- Redistributing I/O costs is not smaller than those of splitting.
- It has a chance that redistributing does not work; thus costs for exploring redistribution are wasted.


# Deleting a Data Entry from a B+ Tree

- Start at root, find leaf  $L$  where entry belongs.
- Remove the entry.
  - If  $L$  is at least half-full, *done!*
  - If  $L$  has only
    - Try to *re-distribute* https://eduassistpro.github.io/ *g* (adjacent node with same parent as  $L$ ).
    - If re-distribution fails, *merge*  $L$
- If merge occurred, must delete entry (pointing to  $L$  or sibling) from parent of  $L$ .
- Merge could propagate to root, decreasing height.

# Deleting 19\*

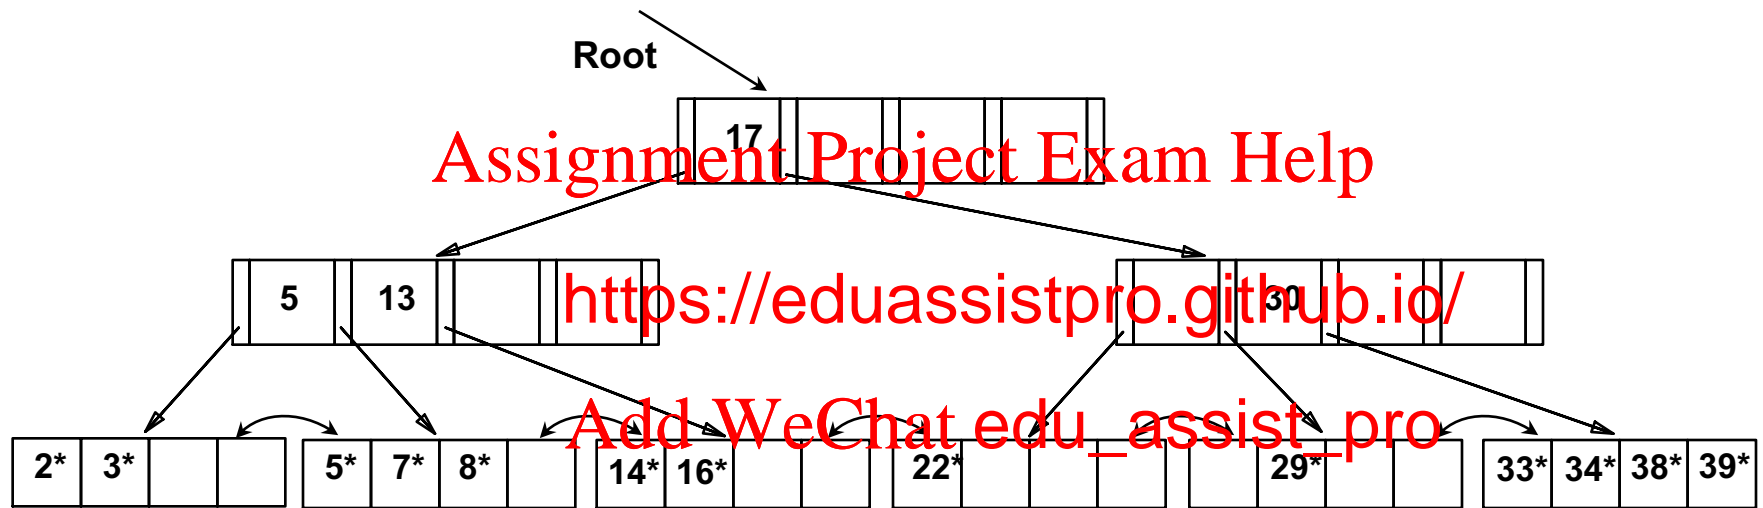


# Deleting 20\*



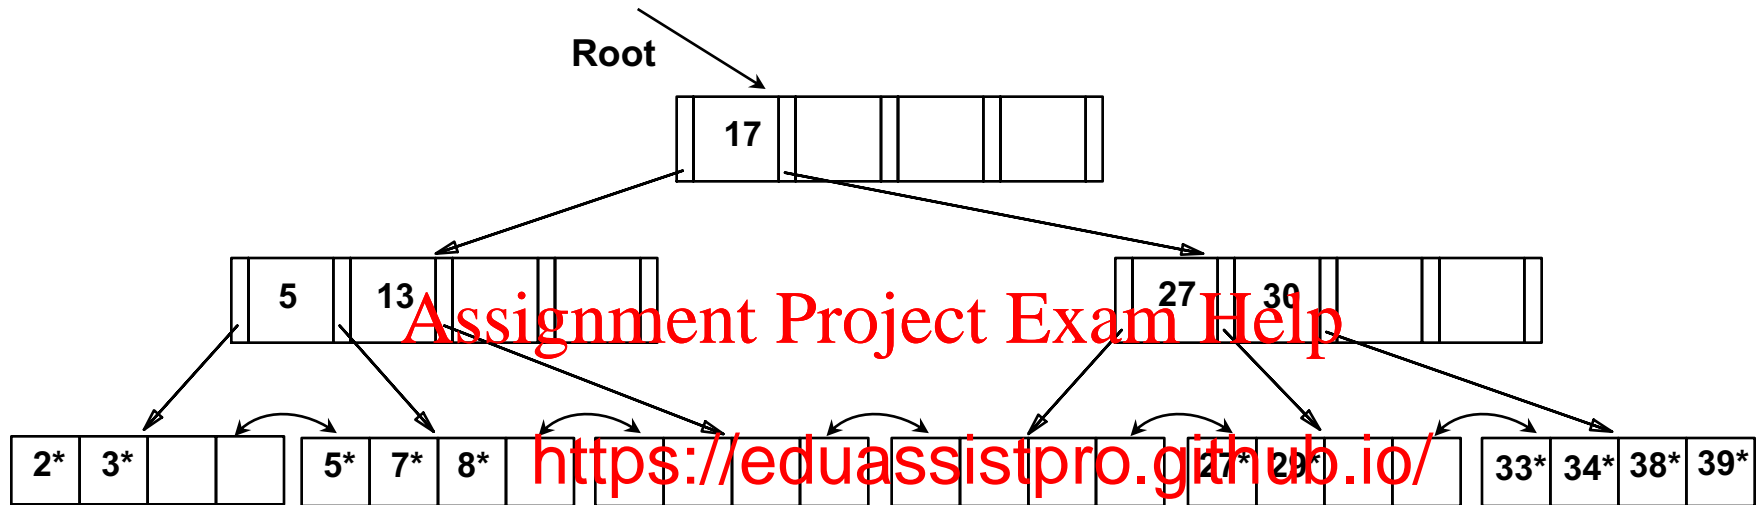
Add WeChat edu\_assist\_pro

# After Deleting 20\* ...




Deleting 20\* is done with re-distribution.  
Notice how middle key is *copied up*.

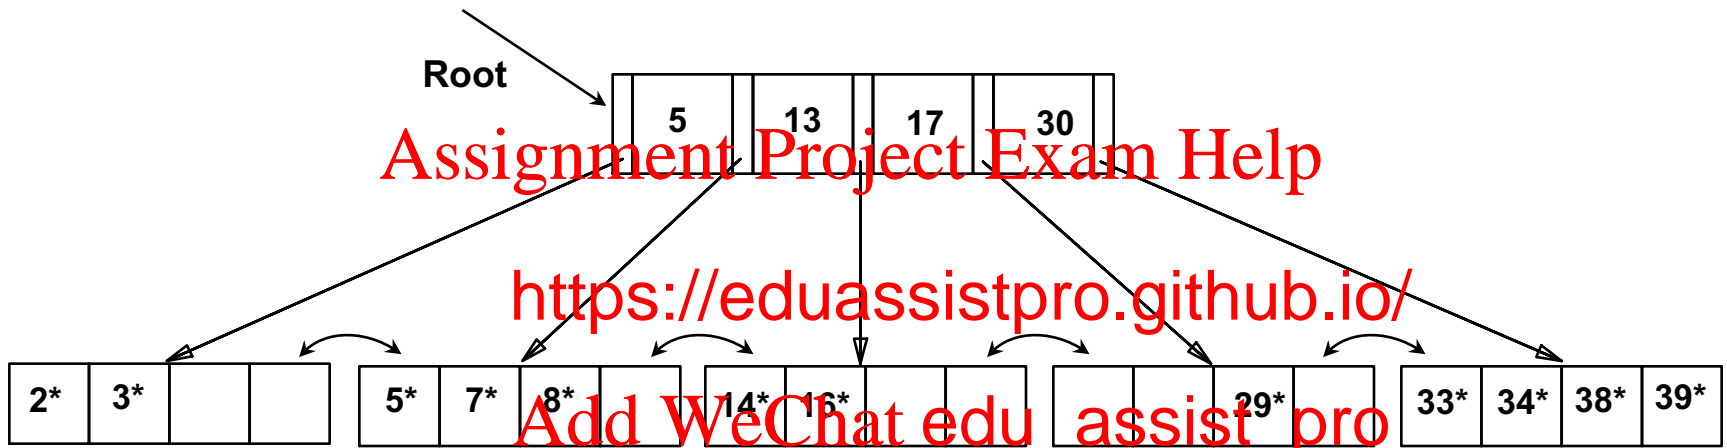
# Deleting 24\* ...



Add WeChat edu\_assist\_pro

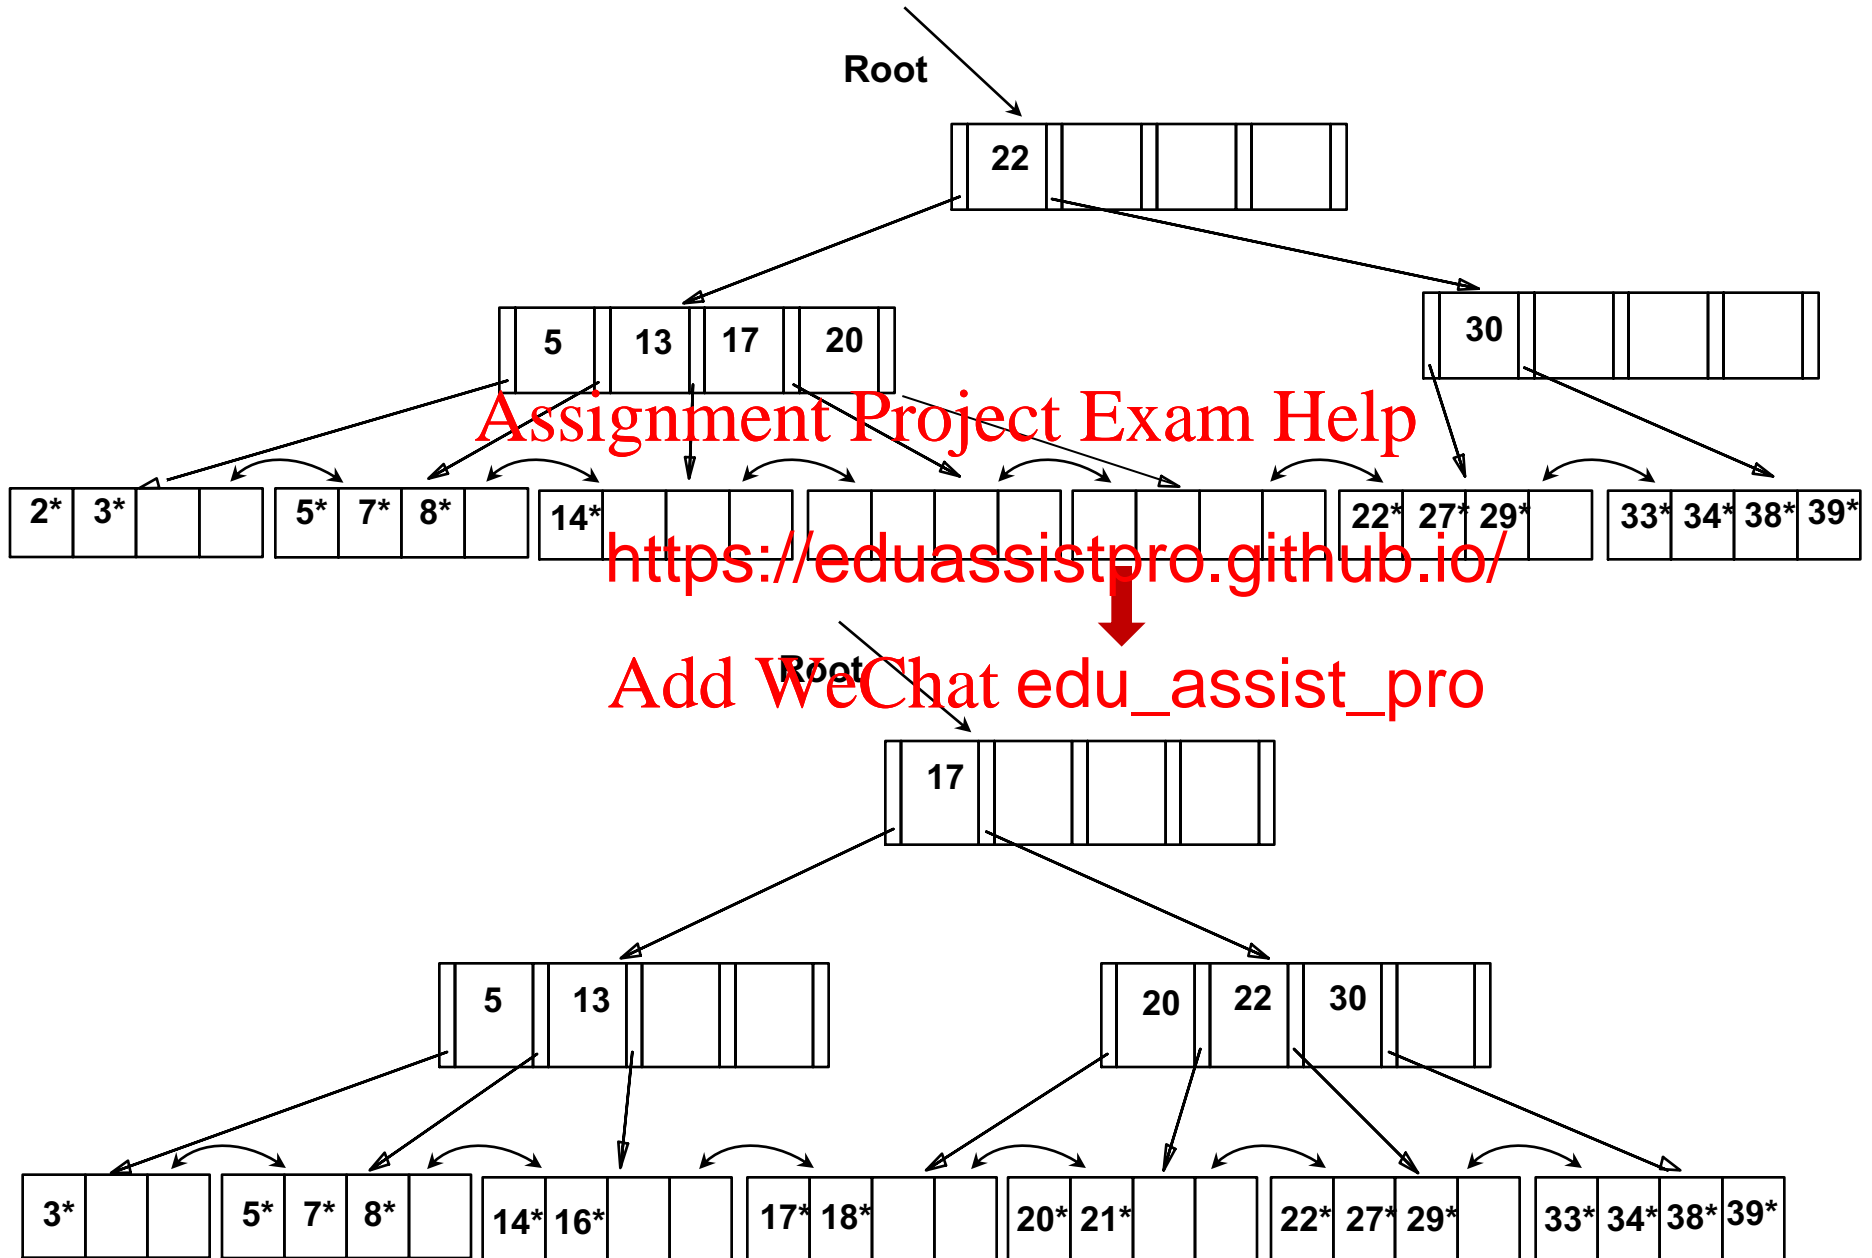


# After Deleting 24\*






# Example of Non-leaf Re-distribution



# Prefix Key Compression

- Important to increase fan-out. (Why?)
- Key values in index entries only 'direct traffic' ; can often compress them.
  - E.g., If we have adjacent index entries with search key values *Dannon Yogurt*, *David Smith*, and *Devarakonda Murthy*, we can abbreviate *David Smith* to *Dav*. (The other keys can be compressed too)
  - Is this correct  
(Can only compress *David Smith* if its index entry greater than every key value (in any subtree) to its left.)
  - In general, while compressing, must ensure its index entry greater than every key value (in any subtree) to its left.
- Insert/delete must be suitably modified.



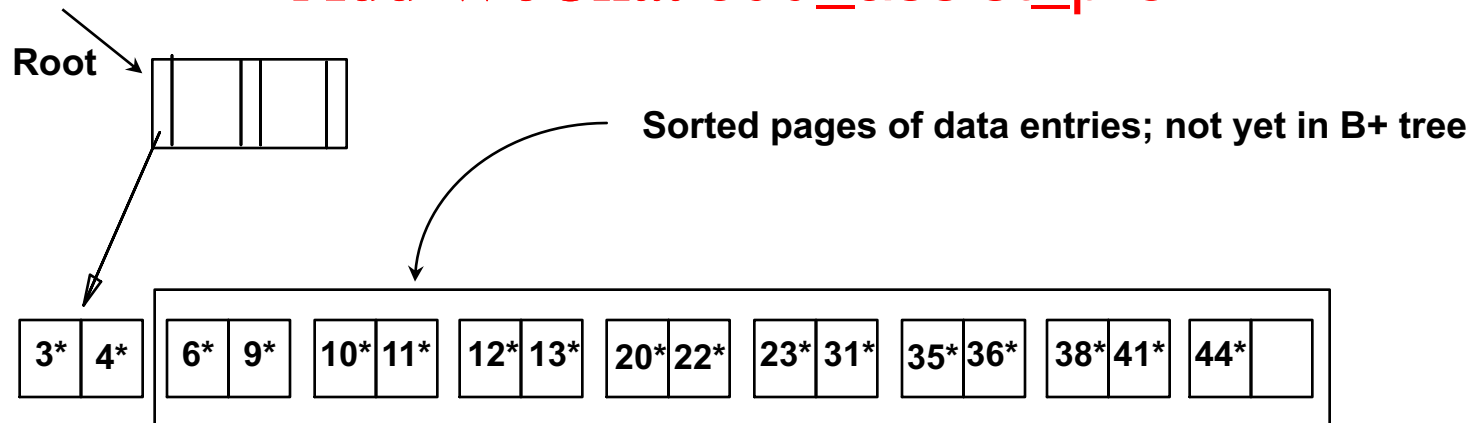
# Bulk Loading of a B+ Tree

- If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
- Bulk Loading can be done much more efficiently.
- *Initialization:* insert pointer to first (leaf) page

Assignment Project Exam Help

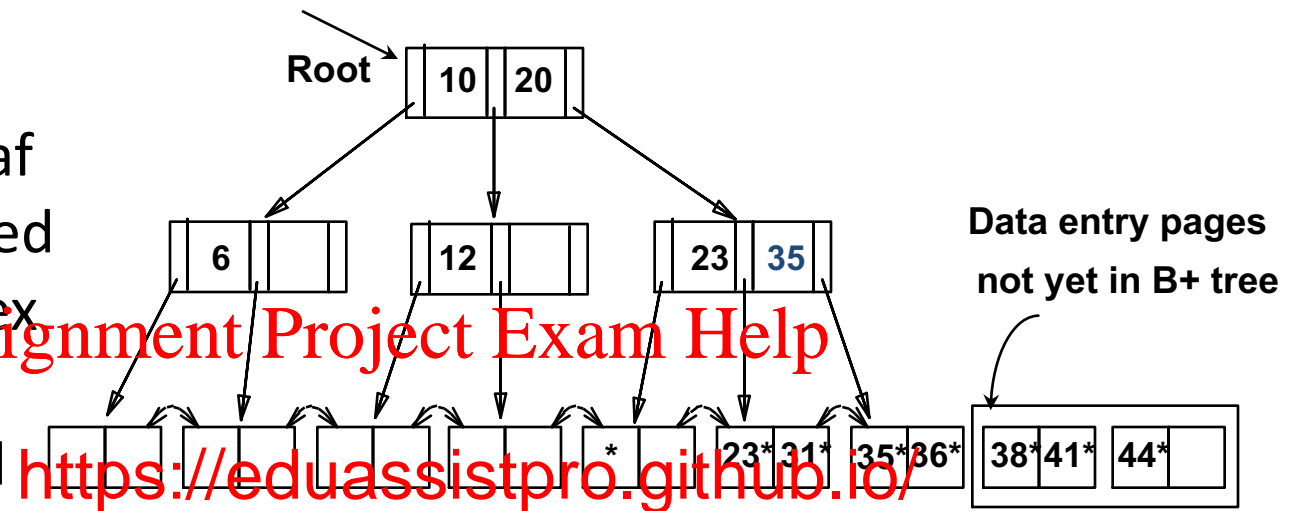
<https://eduassistpro.github.io/>

Add WeChat edu\_assist\_pro




# Bulk Loading (Contd.)

- Index entries for leaf pages always entered into right-most index page just above leaf level. When this fills up, it splits. (Split goes up right-most path to the root.)



<https://eduassistpro.github.io/>



- Much faster than repeated inserts, especially when one considers locking!

# Summary of Bulk Loading

## Option 1: multiple inserts.

- Slow.
- Does not give sequential storage of leaves.

Assignment Project Exam Help

## Option 2: Bulk

- Has advantage of <https://eduassistpro.github.io/> control.
- Fewer I/Os during build.
- Leaves will be stored sequentially (and linked, of course).
- Can control “fill factor” on pages.

Add WeChat edu\_assist\_pro

# A Note on `Order`

*Order* (d) concept replaced by physical space criterion in practice (*`at least half-full`* ).

- Index pages can typically hold many more entries than leaf pages.
- Variable sized keys mean different nodes will contain different sized entries.
- Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries (if we use Alternative (3)).

# Summary

Tree-structured indexes are ideal for range-searches, also good for equality searches.

ISAM is a static structure.

- Only leaf pages needed.
- Overflow chains unless size of data set and data distribution constant.

B+ tree is a dynamic structure

- Inserts/deletes leave tree height-balanced;  $\log_F N$  cost.
- High fanout (**F**) means depth rarely more than 3 or 4.
- Almost always better than maintaining a sorted file.

# Summary (Contd.)

- Typically, 67% occupancy on average.
- Usually preferable to ISAM, modulo locking considerations, adjusts to growth gracefully.
- If data entries can change rid!
- Key compression reduces height.
- Bulk loading can be much faster than repeated inserts for creating a B+ tree on a large data set.
- Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.