

External Sorting

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Why Sort?

- A classic problem in computer science!
- Data requested in sorted order
 - e.g., find students in increasing *gpa* order
- Sorting is first step in *bulk loading* B+ tree index.
- Sorting useful <https://eduassistpro.github.io/> in a collection of records (Why
- *Sort-merge* join algorithm i [Add WeChat edu_assist_pro](https://eduassistpro.github.io/) rting.
- Problem: sort 1Gb of data with 1Mb of RAM.
- What is the minimum number of buffer pages needed to sort a file with arbitrary size?
 - Three.

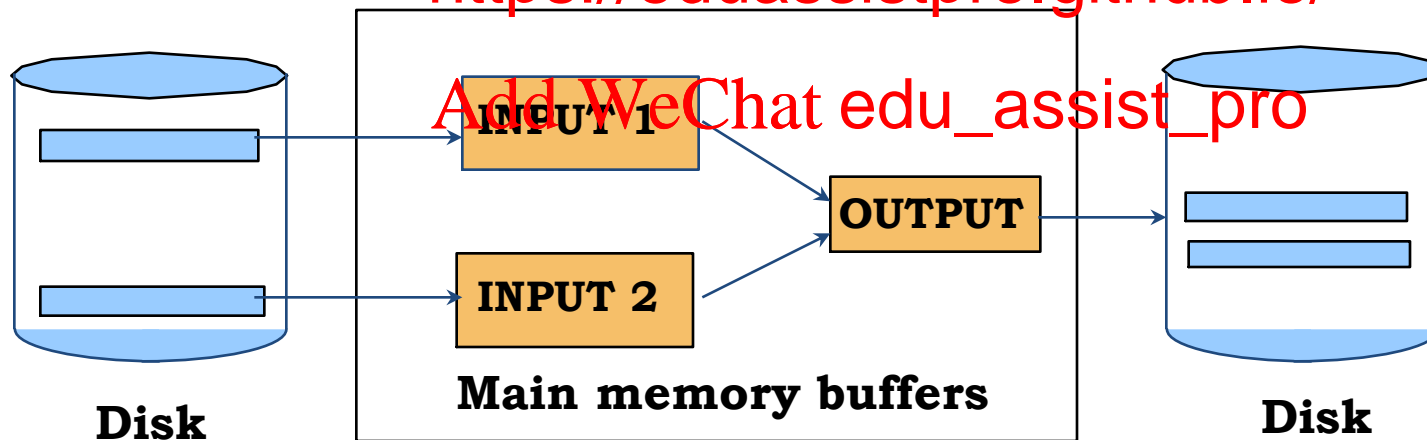
2-Way Sort: Requires 3 Buffers

- Pass 1: Read a page, sort it, write it.
 - only one buffer page is used
- Pass 2, 3, ..., etc.:
 - three buffer

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

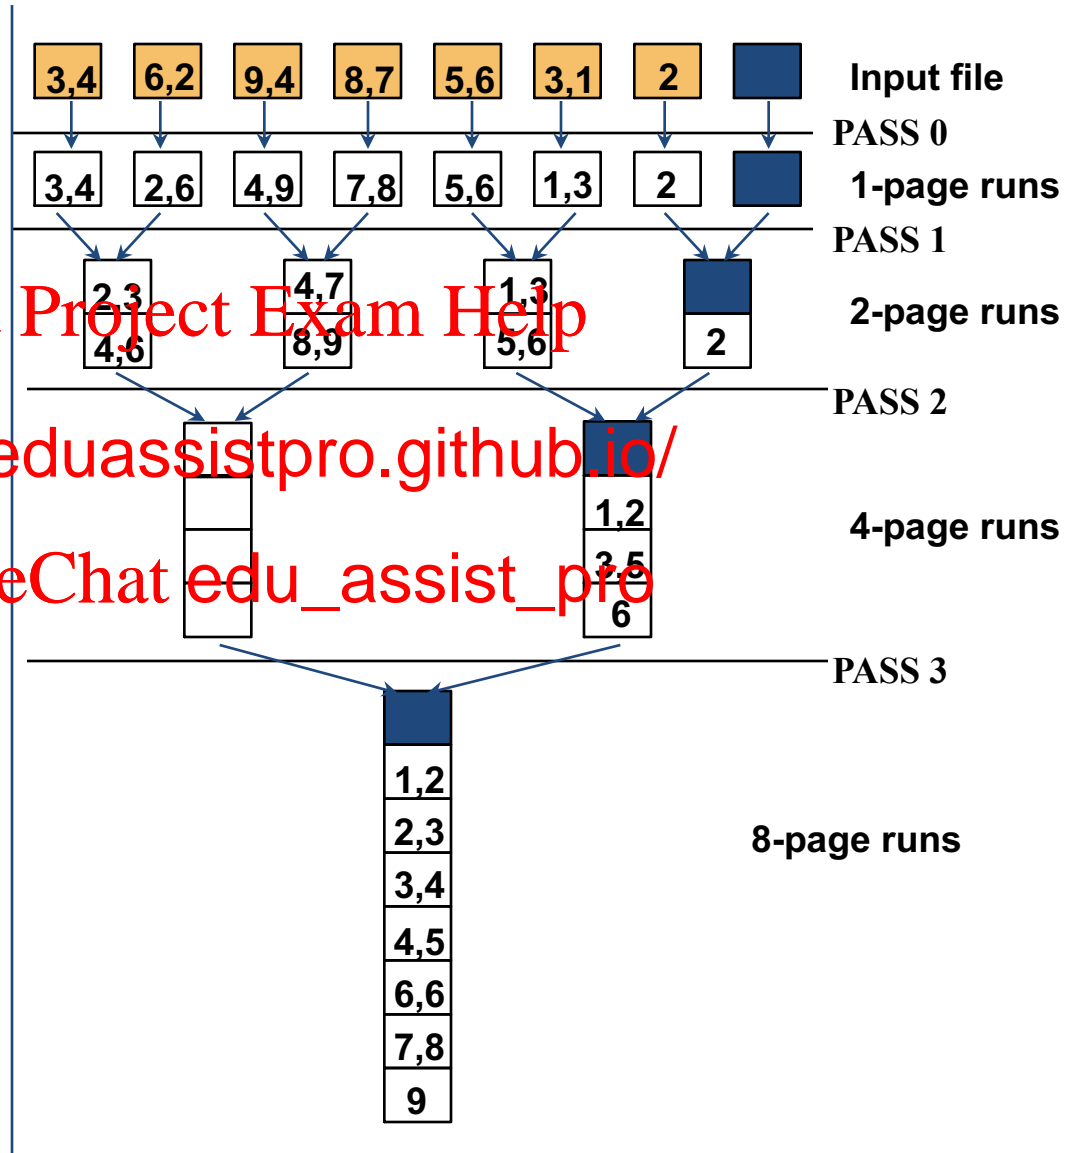


Two-Way External Merge Sort

- Each pass we read + write each page in file.
- N pages in the file => the number of passes

$$= \lceil \log_2 N \rceil + 1$$
- So total cost is:

$$2N(\lceil \log_2 N \rceil + 1)$$
- Idea: **Divide and conquer**:
sort subfiles and merge



General External Merge Sort

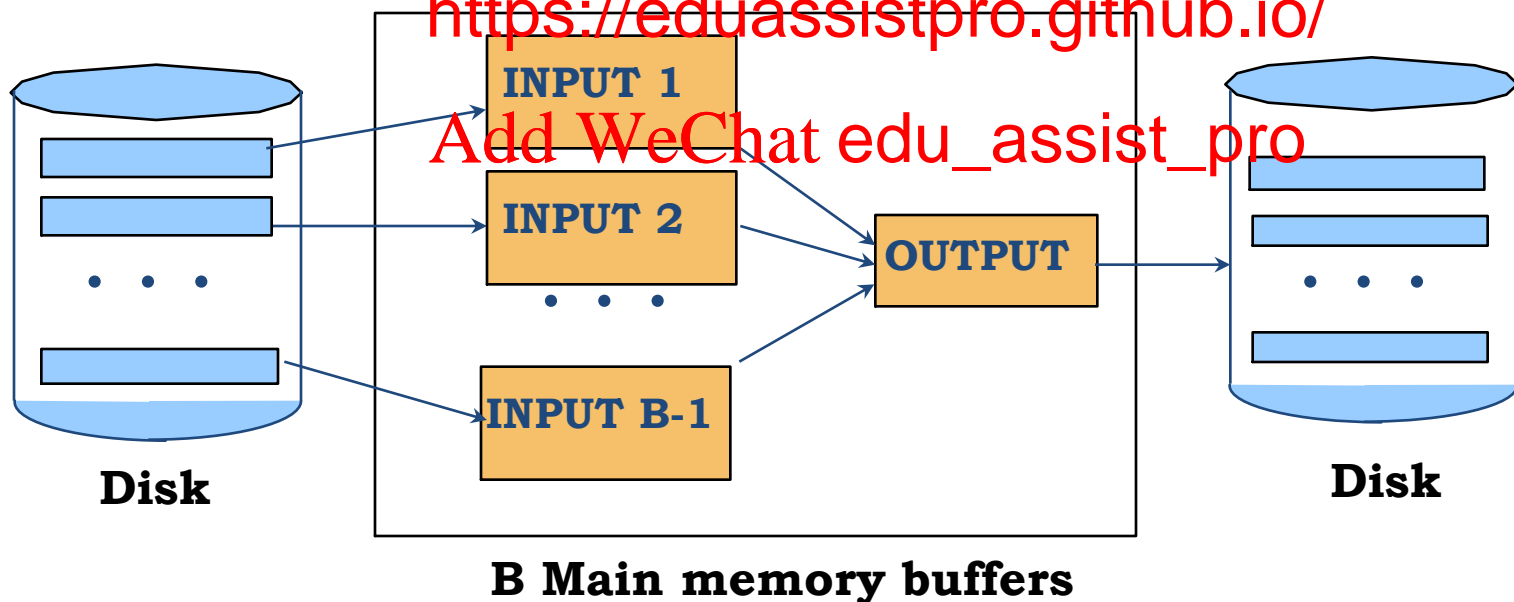
➡ *More than 3 buffer pages. How can we utilize them?*

- To sort a file with N pages using B buffer pages:
 - Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.
 - Pass 2, ..., etc.: merge $B-1$ runs.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



4, 3

6, 2

9, 4

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

4

6

9, 4

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Cost of External Merge Sort

- Number of passes:
- $\text{Cost} = 2N * (\# \text{ of passes})$ $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- E.g., with 5 buffer pages, to sort 108 page file:
 - Pass 0: $\lceil 108 / 5 \rceil = 22$ pages each
(last run is 3 pages)
<https://eduassistpro.github.io/>
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs
(last run is only 3 pages)
[Add WeChat: edu_assist_pro](https://eduassistpro.github.io/)
 - Pass 2: 2 sorted runs, 80 pages and 28 pages
 - Pass 3: Sorted file of 108 pages

Number of Passes of External Sort

N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000					2	2
10,000					2	2
100,000	17	9			3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Speed-up: Internal Sort Algorithm

Quicksort is a fast way to sort in memory.

An alternative is “tournament sort” (a.k.a. “heapsort”): average run length is $2B$.

- **Top:** Read in B blocks
 - **Output:** offer
 - Read in a new
 - insert r into “
 - if r not smallest,
 - then **GOTO Output**
 - else
 - remove r from “heap”
 - output “heap” in order; **GOTO Top (next run)**

This can be only effectively used in the first pass. Virtually, make B be $2B$.

I/O for External Merge Sort

- ... longer runs often means fewer passes!
- Actually, do I/O a page at a time
- In fact, read a *block* of pages sequentially!
- Suggests we use buffers (input/output)
 - But this will reduce fan-out merge passes!
 - In practice, most files still sorted in 2-3 passes.

Assignment Project Exam Help

<https://eduassistpro.github.io/>
es.

Add WeChat edu_assist_pro

Number of Passes of Optimized Sort

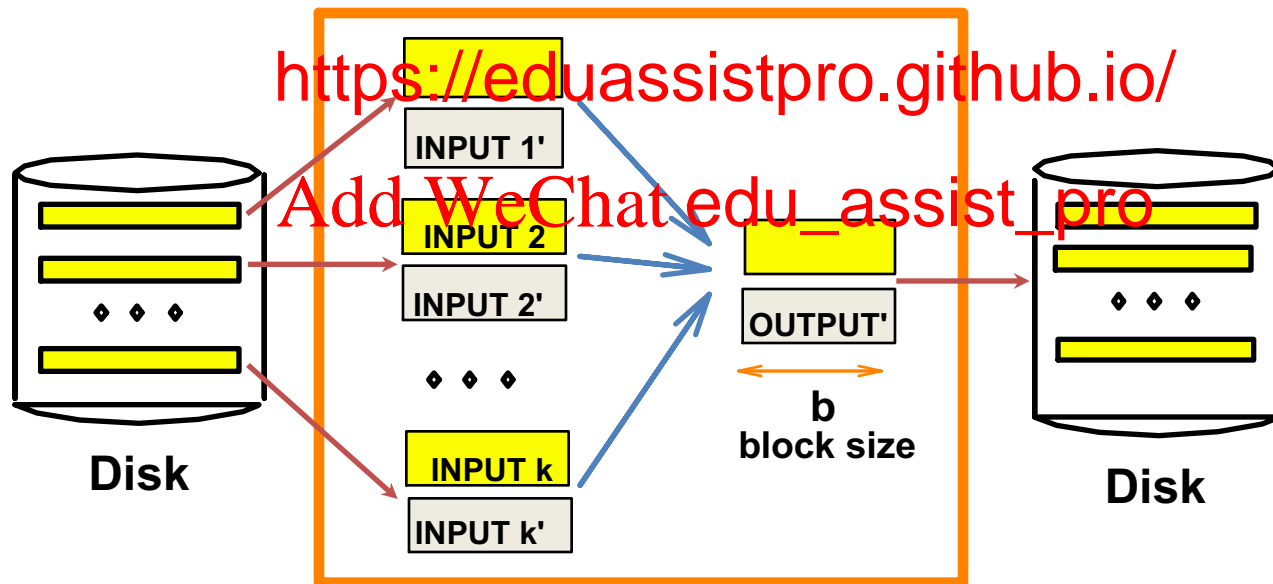
N	B=1,000	B=5,000	B=10,000
100	1	1	1
1,000	1	1	1
10,000	1	1	1
100,000	3	2	2
1,000,000	3	2	2
10,000,000	4	3	3
100,000,000	5	3	3
1,000,000,000	5	4	3

👉 Block size = 32, initial pass produces runs of size $2B$.

Double Buffering

To reduce wait time for I/O request to complete, can *prefetch* into 'shadow block'.

- Potentially, more passes; in practice, most files still sorted in 2-3 passes.



<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Sorting Records!

- Sorting has become a blood sport!
 - Parallel sorting is the name of the game ...
- Datamation: Sort 1M records of size 100 bytes
 - Typical DBMS
 - World record
 - 12-CPU SGI machine, 96 RAM
- New benchmarks proposed:
 - Minute Sort: How many can you sort in 1 minute?

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Using B+ Trees for Sorting

- Scenario: Table to be sorted has B+ tree index on sorting column(s).
- **Idea:** Can retrieve records in order by traversing leaf pages.
- ***Is this a good***
- Cases to consider:
 - B+ tree is clustered ***Good idea!***
 - B+ tree is not clustered ***Could be a very bad idea!***

Assignment Project Exam Help

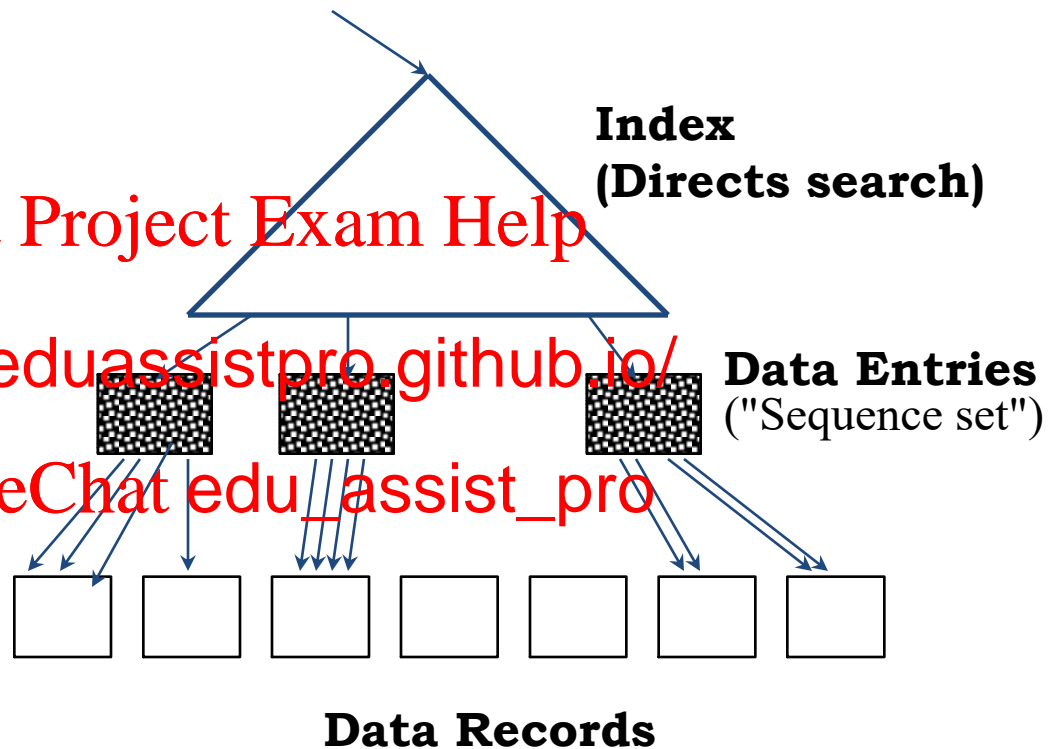
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Clustered B+ Tree Used for Sorting

- Cost: root to the left-most leaf, then retrieve all leaf pages (Alternative 1)

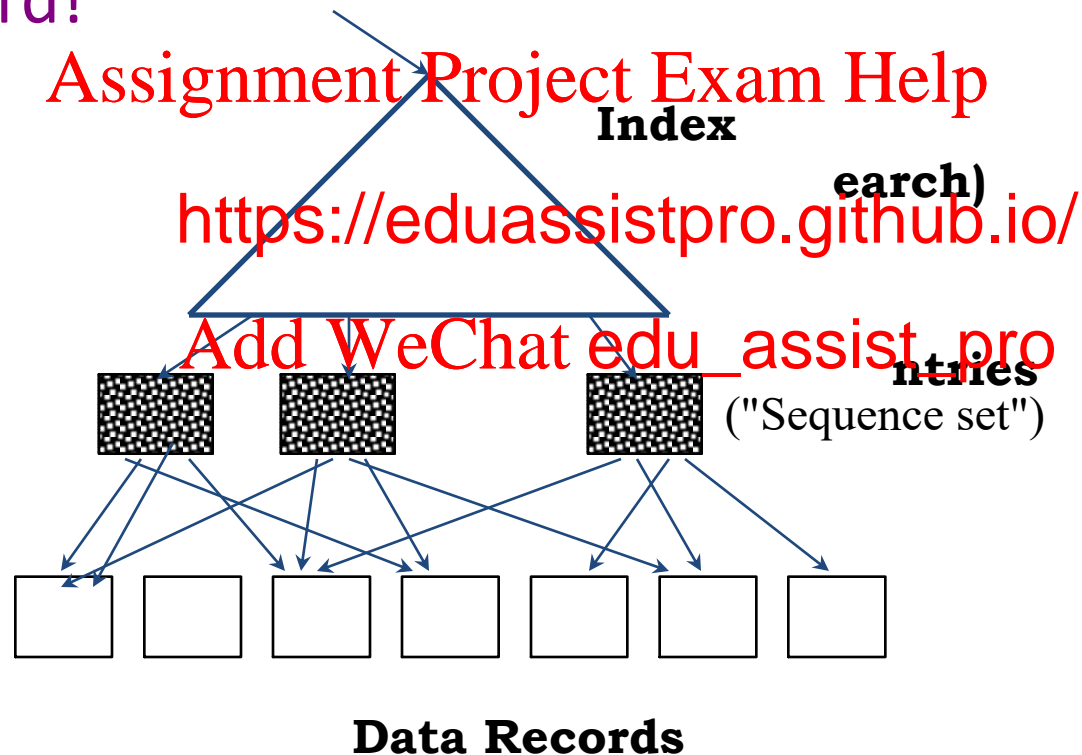
- If Alternative 2 is Additional cost of retrieving data records: each page fetched just once.



➡ *Always better than external sorting!*

Unclustered B+ Tree Used for Sorting

Alternative (2) for data entries; each data entry contains *rid* of a data record. In general, one I/O per data record!



External Sorting vs. Unclustered Index

N	Sorting	p=1	p=10	p=100
100	200	100	1,000	10,000
1,000	2,000			100,000
10,000	40,000		0	1,000,000
100,000	600,000	100,000	10,000,000	10,000,000
1,000,000	8,000,000	1,000,000	10,000,000	100,000,000
10,000,000	80,000,000	10,000,000	100,000,000	1,000,000,000

☞ p : # of records per page

☞ $B=1,000$ and block size=32 for sorting

☞ $p=100$ is the more realistic value.

Summary

- External sorting is important; DBMS may dedicate part of buffer pool for sorting!
- External merge sort minimizes disk I/O cost:
 - Pass 0: Produce B (# buffer pages) runs of size B . Latency is $\frac{N}{B}$.
 - # of runs merged at a time is $\frac{B}{b}$ on B , and **block size**.
 - Larger block size means less I/O cost per page.
 - Larger block size means smaller # runs merged.
 - In practice, # of passes rarely more than 2 or 3.

Summary, cont.

- Choice of internal sort algorithm may matter:
 - Quicksort: Quick!
 - Heap/tournament sort: slower (2x), longer runs
- The best sort
 - Despite 40 years, we're still improving!
- Clustered B+ tree is good for sorting; unclustered tree is usually very bad.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro